

Object Oriented Programming - Exercise 5: Simplified Java Verifier

submission date: 29/1/25

Contents

1	Goals	2
2	Submission Details	2
3	Introduction	2
4	Input/Output	2
5	s-Java specifications	3
5.1	General Description	3
5.2	Variables	4
5.2.1	Variable declaration	5
5.2.2	Assigning values to variables	6
5.2.3	Referring to variables	7
5.3	Methods	7
5.4	if and while Blocks	9
5.5	General Comments	9
6	Important Requirements	10
6.1	Error Handling	10
6.2	Object Oriented Design	11
6.3	Regular Expressions	11
7	Submission Requirements	11
7.1	School Solution	11
7.2	Automatic Testing	12
7.3	README	12
7.4	Submission Guidelines	12

1 Goals

1. Implementing concepts learned in class (Regular Expressions).
2. Designing & implementing a complex system.
3. Working with the Exceptions mechanism.

2 Submission Details

- Submission Deadline: **29/01/2025, 23:55**.
- It is **recommended** to do the exercise in pairs.
- Note: In this exercise you may use the following classes in the standard Java distribution: Classes and interfaces from the `java.util`, `java.text`, `java.io` and `java.lang` packages (including sub-package). Specifically, you are required to use `java.util.regex.Matcher` & `java.util.regex.Pattern` (for working with regular expressions). Also consider `java.util.function` for existing useful functional interfaces.

It is essential to create a UML (Unified Modeling Language) diagram to effectively visualize the system's architecture and design. The UML diagram will help in clarifying the structure, relationships, and behavior of the components, making it easier to communicate ideas and ensure a better understanding across the team. Additionally, it will serve as a valuable reference throughout the development and maintenance phases.

3 Introduction

Regular expressions are used to analyze text by assessing whether a given text string matches a pre-defined pattern. A common setting (which you are now very familiar with) that involves the evaluation of text against a given set of rules is the compilation of programming code. This process involves getting a text file as input, and examining each of its lines to see whether the file is a legal code file, as defined by the language specification.

In this exercise you will implement a Java verifier - a tool able to verify the validity of Java code; it knows how to read Java code and determine its validity, but not to translate it to bytecode. As Java's syntax is rather complicated, writing a Java verifier is a highly challenging task.

To make this task feasible for you, you will in fact implement an *s*-Java verifier (*s* stands for *simplified*), which only supports a **very limited** set of Java features (see details below). Your verifier will not have to run any code, it will only output whether or not the input file is a legal *s*-Java file or not.

4 Input/Output

Your program will receive a single parameter (the *s*-Java source file) and will run as follows:

```
java ex5.main.Sjavac source_file_name
```

This already demands a certain program structure: a package named `ex5`, containing a package named `main`, which in turn contains a class named `Sjavac` which contains the `main`

method of the program. You are free to determine which of the other classes you create will belong in the main package, and which in other packages under the oop.ex5 package. The output of your program is a single digit, outputted using `System.out.println()`:

- 0 – if the code is legal.
- 1 – if the code is illegal.
- 2 – in case of IO errors (see errors).

In case of an error (type 1 or 2), you are also required to print an informative message to the screen using `System.err(error message)`, describing in general what went wrong. There is no specific format you have to follow here, as it won't be tested automatically. However, very uninformative error messages might result in point reduction. See examples below.

5 s-Java specifications

s-Java is a (very) simplified version of Java. In the following, we describe its features.

5.1 General Description

- An s-Java file does not interact with other files. That is, you cannot import code from or export to other files. Each file is stand-alone.
- Each s-Java file is composed of two kinds of components:
 - *Global variables, or members*, are variables shared between methods.
 - *Methods* are general functions. Each method is composed of a list of code lines: defining local variables, giving new values to variables (local or global), calling methods, defining if/while blocks and returning (see more details below).
- s-Java has no classes. As such, it is appropriate to think of s-Java members/global variables as comparable to Java static members, and of s-Java methods as comparable to Java static methods.
- Each s-Java line is either
 - An empty line, containing only white spaces. It should be ignored.
Note: In this exercise, white spaces refer to all characters in the `\s` regex symbol, except for characters that end lines i.e `\n` & `\r`.
However, when you read the input file using the 'readLine' method, each line is a String **that does not contain the end of line characters**. Therefore, **when you analyze a specific line, the `\s` regex can be used to capture white spaces** (it won't capture any end-of-line characters, since they cannot appear in the line in the first place).
 - A comment line, in which the first characters are `//`. No characters, including white spaces, may appear before the `//`, and any number and kind of characters may appear after the `//`, and are to be ignored.

- A code line, which **must** end with one of the following suffixes:
 - * ';' - for defining variables, changing variable values, calling methods and returning.
 - * '{' - for opening method declarations or if/while blocks.
 - * '}' - for closing '{' blocks (has to be in its own line).

White spaces may appear before and after any of these suffixes. These suffixes may **not** appear in the next line instead, as in:

```
void foo()
{ // This is not legal!
```

- Other comment styles like multi-line comments (`/* . . . */`), javadoc comments (`/** . . . */`) and single-line comments appearing in the middle of a line are not supported by s-Java . Any appearance of such comments is illegal and should result in a printed value of 1.
- Operators are not supported (`int a = 3 - 5;`, `String b = "OO" + "P";` are illegal).
- Similarly, arrays are not supported in s-Java .

5.2 Variables

s-Java comes with a strict set of variable types. It does not support the creation of new types (e.g., classes, interfaces, enums). The language supports the definition of two kinds of variables: global variables and local variables (defined inside a method – see below). Both types of variables are defined in the same way as in java:

type name = value;

- Table 1 shows the legal s-Java types.
- *name* is any sequence (length > 0) of letters (uppercase or lowercase), digits and the underscore character ('_'). *name* may **not** start with a digit. *name* may start with an underscore, but in such a case it must contain at least one more character (i.e., '_' is not a legal name). Furthermore, a name may not start with two or more consecutive underscores.

Legal name examples: 'g2', 'b_3', '_a', '_0', 'a_'.

Illegal name examples: '2g', '_', '2_', '54_a', '3_3_3_b', '___', '___b'..

- *value* can be one of the following:
 - a legal value for *type* (see table 1).
 - * Note that both int and double values may be padded with zeros, and may receive + and - signs. Also, double values may start or end with a period, but must contain at least one digit (i.e, double d = .; should not work).
 - * you may assume the following **four** characters will not appear in String or char values:
 \ , ' , " , ,
 - another existing and initialized variable of the *same type*

5.2.1 Variable declaration

- A variable may be declared with or without an initialization. That is, both `int a;` and `int b2 = 5;` are legal s-Java lines.
- A variable declaration must be encapsulated in a single line. For example, the following are

illegal:

```
int a
= 5;
int g
;
```

Type	Description	Value Format	Examples
int	an integer number	a number (positive, 0 or negative)	int num1 = 5; int num2 = -3; int num3 = +001;
double	a real number	an integer or a floating-point number (positive, 0 or negative)	double b = 5.21; double c = 2.; double e = .1; double k = -.3;
String	a string of characters	a string of characters	String s = "hi"; String a = "i%#";
boolean	a boolean variable	true, false or any number (int or double)	boolean a = true; boolean b = 5.2;
char	a single character	any character (inside single quotation marks)	char my cr = 'a'; char g = '@';

Table 1: s-Java types.

- Multiple variables **of the same type** may be declared in a single line, separated by a comma. For example, the following lines are **legal**:

```
double a, b;  
int i1, i2 = 6;  
char c='Z', f;  
boolean a, b, c, d = true, e, f = 5;  
String a = "hello", b = "goodbye";
```

- There may not be two global variables with the same name (regardless of their types). For example, in the following, given the first line, the second line is **illegal**:

```
int a = 5;  
String a = "hello";
```

However, a local variable can be defined with the same name as a global one (regardless of their types). That is, in the previous example, had the first line been a declaration of a global, the second would have been legal if it was a declaration of a local variable. This is true even if the

global variable of the same name was initialized inside the method where the local variable is later declared.

- Methods parameters are considered local variables of the method they belong to.
- Two local variables with the same name (regardless of their type) cannot be defined inside the same block. This also applies to variables with the same name as a method parameter.
- However, two local variable with the same name can be defined inside different blocks, even if one is nested in the other.

```
void foo(int param1){
    int a = 5;
    if (param1){
        int a = 20;
        boolean param1 = true;
        while (param1) {
            double a = 2.5;
        }
    }
    return;
}
```

In case of multiple variables with the same name, a reference to that name refers to the variable in the most specific scope (or, in case of global and local variables, to the local variable).

- A variable (local or global) may have the same name as a method.

5.2.2 Assigning values to variables

- Any variable may be declared using the modifier **final**, which makes it a constant. Such variables **must** be initialized with some value at declaration time: **final int a = 5;**. The modifier should appear before the type of the variable, and not after. In other words, the following is not allowed: **int final a = 5;**.
- A final variable may **not** be assigned a value in subsequent lines (i.e., lines other than it's declaration line).
- Using the final modifier on a line with more than one initialization makes all of the variables final. For example:

```
int final a = 521, b = 9; → both a and b are final.
```

- Other java modifiers (e.g., static, public/private) are **not** allowed (they are not part of the language specification).
- A (non final) variable can be assigned with a value after it is created. **It should be allowed to assign multiple variables in the same line:**

```
int a = 521;
String b;
...
a = 832;
a = 1, b = "hi";
```

This applies both to global and local variables. Local variables can only be assigned inside the method they were declared in, including any scopes nested in it (since they can only be accessed inside the scope they were declared in - see section 5.2.3). Global variables may be assigned multiple times both inside and outside a method.

- A local variable cannot be used in any way (in an assignment, in a call to a method) before it is assigned a value (you don't have to check that the variable initialization is reachable, e.g., inside an if (false) block). In other words, if a local variable is not initialized in its declaration, the code is only legal if:

The next line **in which it appears** is an assignment of a value to that variable.

It is never used.

- A variable a (global or local) can be assigned with another (global or local) initialized variable b of the same type (it could be on the same line). For example:

int a = 5, b ;

b = 7, a=b;

- Additionally, a double can also be assigned with an int, and a boolean can also be assigned with an int and a double. For example:

int a = 5;

double b = a

5.2.3 Referring to variables

You cannot access a variable that is defined in a scope more internal than the scope you are currently in. Access to a variable is possible under one of the following conditions:

- It's a local variable that has been defined in the lines preceding it in our scope (or in a more external scope).
- It's a global variable that was defined somewhere in the external scope, and we are inside any method.
- It's a global variable, and we are in the external scope after the line where it was defined.

When inside a method and there's a variable accessible to us (unless it's a global variable initialized in the external scope), the question of whether it's initialized or not depends on if its initialization took place in any of the preceding lines of the method, even if it's within a more internal scope. This applies to both global and local variables.

5.3 Methods

s-Java methods are a simple version of Java's methods definition:

```
void method name ( type1 parameter1, type2 parameter2, . . . , typen parametern) {
```

where:

- *method name* is defined similarly to variable names (i.e., a sequence of length > 0, containing letters (uppercase or lowercase), digits and underscore), with the exception that method names must start with a letter (i.e., they may not start with a digit or an underscore).
- *parameters* is a comma-separated list of parameters. Each *parameter* is a pair of a valid type and a valid variable name, **without** a value.
- Only void methods are supported.

After the method declaration, comes the method's code. It may contain the following lines:

- Local variable declaration lines (as defined in Section 5.2)
- Variable assignment lines (as defined in Section 5.2).
- A call to another existing method. Any method `foo()` may be called, **regardless** of its location in the file (i.e., before or after the definition of the current method). The syntax of calling is the java syntax:

```
method name (param1, param2, . . . , paramn);
```

where *method name* is an existing method and *param*₁, *param*₂, . . . , *param*_n are variables or constants (3, "hello") of types agreeing with the method definition (though a double parameter can accept an int argument and a boolean can int and double). Calling a method with an incompatible number of arguments, wrong types or uninitialized variables is **illegal**.

- An if/while block (see below).
- A return statement. Return statements may appear anywhere inside a method, but must also appear as the last line in the method's code (you do not need to check for unreachable code due to previous return statements); this is different from Java where a method can also simply end without a return statement. The return statement should not return any value, in accordance with the method's declaration. The format is:

```
return;
```

Note the following:

- Method parameters may be final. For example:

```
void foo(final int a) {
```

In this case, it is still legal to call the method with a non-final variable. It only means that the value of the parameter may not be changed inside the method.

- You may change the value of a non-final parameter inside the method's code.

- Methods must end with a line containing the single token ' ' (which comes right after the return line as presented above).
- Recursive calls are allowed. I.e., a method may call itself. You are not required to check for infinite loops/recursions in the code.
- Method overloading is not supported: no two methods with the same name may exist.
- The order in which methods appear in the code is meaningless. Methods may also appear before/after/between global variable declarations.
- A method may **not** be declared inside another method.

5.4 if and while Blocks

if and while blocks may only appear **in methods**, and are defined in the following way:

```
if (condition) {  
    ...  
}  
while (condition) {  
    ...  
}
```

where condition is a boolean value (as defined in table 1), so it is either:

- One of the reserved words is true or false.
- An initialized boolean, double or int variable.
- A double or int constant/value (e.g. 5, -3, -21.5).

Also, multiple conditions separated by AND/OR may appear (e.g., if (a || b || c) { is a legal statement).

Notice the following:

- The || and && operators must be placed between two boolean expressions, so for example if(|| a) is illegal, and so is if(a || && b).
- if/while blocks may contain the same type of lines as methods (that is, variable declaration, method calls, etc. See Section 5.3).
- You are not required to support conditions containing brackets, like

```
if (( a || b) && c ...) {  
    ...  
}
```

- Much like methods, if/while blocks must start with a '{' token (which should be at the end of the condition line) and end with a '}' token (which should be in a line of its own).
- if/while blocks can be nested to a practically unlimited depth (i.e. you should support a depth of at least java.lang.Integer.MAX VALUE): if inside while and vice versa.
- You are not required to support else if, else blocks, do/while loops, for loops or switch statements. That is, the words else, do, for and switch are not part of the s-Java specification.

5.5 General Comments

- A main method is not required in s-Java . Of course a method called main may be defined, much like any other method.
- Naming variables or methods with reserved words (e.g int, if...) is not covered by the s- Java specification, and so you can handle such cases as you wish. (For example: int double = 6;)

- White spaces are allowed anywhere in the code, and are to be ignored. This **excludes** white spaces inside variable names, values or reserved words (e.g., the following line is **illegal**: `i n t a = 5;`), and in the start of comment lines, so the following line is illegal:
`" \comment after spaces"`

- On the other hand, white spaces are required to separate between:
 - A method's return type (only void in s-Java) and name (e.g., `voidfoo(){}` is illegal).
 - A variable's type and name (e.g., `inta;` is **illegal**).
 - A final modifier and a variable type (e.g., `finalint a;` is **illegal**).

Zero or more white spaces may appear between any other input tokens (this includes before/after parentheses, '=', ',', '{' and '}' characters, etc.).

- Each line of code must appear in a single line, and not broken into several lines. Also, no two statements may appear in the same line.
- Packages, as well as exceptions, are not supported in s-Java . That is, the words package, try, catch and finally are not part of the s-Java specification.
- The reserved keywords in s-Java are:
 - Types: int, double, boolean, char, String.
 - Other: void, final, if, while, true, false, return.

6 Important Requirements

6.1 Error Handling

You are required to use the exceptions mechanism to handle general errors of your program (e.g., an IOException caused by an illegal file name). As noted earlier, in such cases you should print the value 2 and an understandable error message.

When using exceptions, note that:

- Exception class should be called XXXException, and should be put in the same package as the class that throws it (not in a separate 'exceptions package'). It should not be defined as a nested class.
- Be sure to catch the most specific kind of exception that is relevant in the context - i.e, if your method throws a specific kind of exception, catch this specific exception and not a more general one.

Regarding the main task (deciding whether or not the file is a legal code file), error handling is a bit tricky. Supposedly, all the information many of your methods have to provide in this exercise is a single bit – whether or not some part of the code is legal. In this case, it is tempting to define boolean methods that return true or false according to whether the code is legal or not. However, some of the benefits of the exceptions mechanism might come in handy here. We advise you to consider using this mechanism in this exercise. You're advised to implement specific exception classes for different exceptions.

6.1.1 IO Exceptions

The following should raise an IO Exception:

- Illegal number of arguments for the program
- Wrong file format (not sjava)
- Invalid file name

6.2 Object Oriented Design

As in previous exercises in this course, your program should follow the object oriented design principles. While the task here is very different than in previous exercises, the working process should be similar:

- Try to divide your program into small, independent units.
- Try to think of several design alternatives for each unit, consider the pros and cons of each alternative, and select the one you think is best.

You are required to create a UML (Unified Modeling Language) diagram to effectively visualize the system's architecture and design. The UML diagram will help in clarifying the structure, relationships, and behavior of the components, making it easier to communicate ideas and ensure a better understanding across the team. Additionally, it will serve as a valuable reference throughout the development and maintenance phases.

6.3 Regular Expressions

One of the main goals of this exercise is to practice the use of the regular expression mechanism in Java (and in general). You are therefore required to make extensive use of it in your program. However, you are allowed (and encouraged!) to use other text analyzing mechanisms in your code (e.g., String class methods such as `substring()`, `charAt()`, etc.). The general rule of thumb is that you should use regular expressions whenever it makes your life easier, and not when it makes it harder.

You are encouraged to use online tools to test your regular expressions (i.e regex101), Please notice the special escaping of java, for example a digit in regex101 will be denoted by ' `d` ' while in Java it will be '`\\d`'

Use the tools learned in class (such as boundaries, different types of quantifiers, etc.) to make your regular expressions as efficient and clear as possible.

In your README file, please describe the two main regular expressions you used in your code. For each regular expression, include an explanation of its structure, detailing the purpose of each component within the regex and how it contributes to the overall pattern.

7 Submission Requirements

7.1 School Solution

The school solution may be found under `oop/bin/ex5`. You are strongly advised to experiment with it before starting to work on your design, in order to get a feeling of how your program should

behave on different inputs.

How to run the school solution?

- 1) On the aquarium computers save the sjava file, which you want to execute the solution on, in any directory that is not your home directory (e.g., on the Desktop or in a test directory you create yourself).
- 2) Make sure that your sjava file has reading-permissions: use the 'chmod 755' command on the directory containing the sjava file and also on the sjava file itself.
- 3) Navigate into the directory where the file is located!
- 4) From within the directory, run the command: `oop/bin/ex5/ex5_school_solution file.sjava`. Please note not to do this from outside the directory where the file is located (by writing the full path of the file) but do it from within the directory and just write its name (for example, `file.sjava`).

Note that the school solution does not provide informative comments on the type of error, although you must implement them as instructed.

7.2 Automatic Testing

A file called supplied material.zip can be found in the course website. The zip contains a tests folder with the presubmission tests. The folder includes a description of the presubmission tests on "presubmission_sjavac_tests.txt".

The presubmission will produce error messages containing the file number of failed tests. For example, you may receive the following error:
boolean problem in test number:011 expected:<[0]> but was:<[1]>

This means that your program printed a different value than the school solution (1 instead of 0), when tested with file test011.sjava. This file can be found in the tests sub-folder of the supplied material. To get a short description of that test, go to file presubmission sjavac tests.txt (also supplied), and search for test number 011. Following the name of the file and the expected return value (0) you will find the description: "boolean member test no value". I.e, the test checks how your verifier handles declaration of a boolean member variable with no assigned value.

You can also run the presubmission script directly, by downloading the supplied material.zip folder, inserting your jar file into the exercise jar sub-folder and running the python script run presubmission.py.

However, the presubmission tests are less than 10% of the s-Java files on which your program will be automatically tested. **You are therefore encouraged to test your program on your own.**

7.3 README

Follow the README instructions from the 'Programming Exercise Preparation and Submission' file that appears on moodle. If you are submitting individually, enter only your own CS username on the first line of the README. To submit as a pair, the first two lines of the README must be your CS usernames, and only the student whose CS username is the first line of the README should submit the exercise. When you submit the exercise, make sure that both your usernames appear in the output of the presubmission tests.

In your README file, please describe the two main regular expressions you used in your code. For each regular expression, include an explanation of its structure, detailing the purpose of each component within the regex and how it contributes to the overall pattern.

7.4 Submission Guidelines

You should submit a file named `ex5.jar/ex5.zip/ex5.tar` containing all `.java` files of your program, as well as your README file, and a directory named `tests` with 20 `s-Java` files on which you tested your program.

Java files should be submitted in the directories of their original packages. No `.class` files should be submitted. Make sure that javadoc accepts your program and produces documentation.

Also, Remember that your program must compile without any errors or warnings. To compile your code in command line with the correct configuration, type:

```
javac -Xlint:rawtypes -Xlint:empty -Xlint:divzero -Xlint:deprecation
```

In order to have the warnings displayed when you compile your code in IDEA, go to

File->Settings->Build, Execution, Deployment-> Compiler>Java Compiler and enter the following line under Additional command line parameters:

```
-Xlint:rawtypes -Xlint:empty -Xlint:divzero -Xlint:deprecation
```

Good Luck!